

### **1. simple vector addition:**

```
import tensorflow as tf

# Define two vectors

vector1 = tf.constant([1, 2, 3], dtype=tf.float32)
vector2 = tf.constant([4, 5, 6], dtype=tf.float32)

# Perform vector addition

result = tf.add(vector1, vector2)
```

### **2. Implement a regression model in Keras.**

```
import tensorflow as tf

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

import numpy as np

# Generate some example data

np.random.seed(0)
x_train = np.random.rand(100, 1)
y_train = 2 * x_train + 1 + 0.1 * np.random.randn(100, 1)

# Define the model

model = Sequential([
    Dense(1, input_dim=1)
])

# Compile the model

model.compile(optimizer='adam', loss='mse')

# Train the model

model.fit(x_train, y_train, epochs=100, batch_size=10)

# Make some predictions
```

```
x_test = np.array([[0.5], [1.5]])
predictions = model.predict(x_test)
```

```
# Print predictions
```

```
print("Predictions:", predictions)
```

### **3. Implement a perceptron in TensorFlow/Keras Environment.**

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
import numpy as np
```

```
# Generate some example data
```

```
np.random.seed(0)
```

```
x_train = np.random.rand(100, 2) # 100 points with 2 features
```

```
y_train = (x_train[:, 0] + x_train[:, 1] > 1).astype(int) # Label 1 if the sum of the features > 1, else 0
```

```
# Define the model
```

```
model = Sequential([
    Dense(1, input_dim=2, activation='sigmoid') # Single perceptron unit
])
```

```
# Compile the model
```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
# Train the model
```

```
model.fit(x_train, y_train, epochs=100, batch_size=10)
```

```
# Make some predictions
```

```
x_test = np.array([[0.1, 0.9], [0.8, 0.2]])
```

```
predictions = model.predict(x_test)
```

```
# Print predictions
```

```
print("Predictions:", predictions)
```

### **4. Implement a Feed-Forward Network in TensorFlow/Keras.**

```

import tensorflow as tf

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dense

import numpy as np

# Generate some example data

np.random.seed(0)

x_train = np.random.rand(100, 1) # 100 points with 1 feature

y_train = 2 * x_train + 1 + 0.1 * np.random.randn(100, 1) # Linear relationship with some noise

# Define the model

model = Sequential([

    Dense(10, input_dim=1, activation='relu'), # First hidden layer with 10 neurons and ReLU activation

    Dense(10, activation='relu'), # Second hidden layer with 10 neurons and ReLU activation

    Dense(1) # Output layer with 1 neuron (for regression)

])

# Compile the model

model.compile(optimizer='adam', loss='mse')


# Train the model

model.fit(x_train, y_train, epochs=100, batch_size=10)


# Make some predictions

x_test = np.array([[0.5], [1.5]])

predictions = model.predict(x_test)

```

```

# Print predictions

```

```

print("Predictions:", predictions)

```

## **5. Implement an Image Classifier using CNN in TensorFlow/Keras.**

```

import tensorflow as tf

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dense, Conv2D, MaxPooling2D, Flatten

```

```

from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical

# Load the MNIST dataset

(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Reshape the data to include a channel dimension

x_train = x_train.reshape(-1, 28, 28, 1).astype('float32') / 255.0
x_test = x_test.reshape(-1, 28, 28, 1).astype('float32') / 255.0

# One-hot encode the labels

y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

# Define the model

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

# Compile the model

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model

model.fit(x_train, y_train, epochs=10, batch_size=64, validation_split=0.1)

# Evaluate the model

test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test accuracy: {test_acc:.4f}")

```

## 6. Improve the Deep learning model by fine tuning hyper parameters

```
import tensorflow as tf

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dense, Conv2D, MaxPooling2D, Flatten, Dropout

from tensorflow.keras.datasets import mnist

from tensorflow.keras.utils import to_categorical


# Load the MNIST dataset

(x_train, y_train), (x_test, y_test) = mnist.load_data()


# Reshape the data to include a channel dimension

x_train = x_train.reshape(-1, 28, 28, 1).astype('float32') / 255.0
x_test = x_test.reshape(-1, 28, 28, 1).astype('float32') / 255.0


# One-hot encode the labels

y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)


# Define the model

model = Sequential([

    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),

    MaxPooling2D((2, 2)),

    Dropout(0.25), # Adding dropout to reduce overfitting

    Conv2D(64, (3, 3), activation='relu'),

    MaxPooling2D((2, 2)),

    Dropout(0.25), # Adding dropout to reduce overfitting

    Flatten(),

    Dense(128, activation='relu'),

    Dropout(0.5), # Adding dropout to reduce overfitting

    Dense(10, activation='softmax')

])
```

*# Compile the model with a custom learning rate*

```
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)
```

```
model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
```

*# Train the model*

```
model.fit(x_train, y_train, epochs=20, batch_size=32, validation_split=0.1)
```

*# Evaluate the model*

```
test_loss, test_acc = model.evaluate(x_test, y_test)
```

```
print(f"Test accuracy: {test_acc:.4f}")
```

## **7. . Implement a Transfer Learning concept in Image Classification.**

```
import tensorflow as tf
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Dense, Flatten, Dropout
```

```
from tensorflow.keras.applications import VGG16
```

```
from tensorflow.keras.datasets import cifar10
```

```
from tensorflow.keras.utils import to_categorical
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

*# Load the CIFAR-10 dataset*

```
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

*# Normalize the pixel values*

```
x_train = x_train.astype('float32') / 255.0
```

```
x_test = x_test.astype('float32') / 255.0
```

*# One-hot encode the labels*

```
y_train = to_categorical(y_train, 10)
```

```
y_test = to_categorical(y_test, 10)
```

*# Load the VGG16 model pre-trained on ImageNet, excluding the top fully connected layers*

```
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
```

```
# Freeze the layers of the base model

for layer in base_model.layers:

    layer.trainable = False


# Define the model
model = Sequential([

    base_model,

    Flatten(),

    Dense(256, activation='relu'),

    Dropout(0.5),

    Dense(10, activation='softmax')

])


# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])


# Data augmentation
datagen = ImageDataGenerator(

    rotation_range=15,

    width_shift_range=0.1,

    height_shift_range=0.1,

    horizontal_flip=True

)

datagen.fit(x_train)


# Train the model
model.fit(datagen.flow(x_train, y_train, batch_size=32), epochs=20, validation_data=(x_test, y_test))


# Evaluate the model
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test accuracy: {test_acc:.4f}")
```

## 8.Using a pre trained model on Keras for Transfer Learning

```
import tensorflow as tf

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Dropout

from tensorflow.keras.applications import MobileNetV2

from tensorflow.keras.datasets import cifar10

from tensorflow.keras.utils import to_categorical

from tensorflow.keras.preprocessing.image import ImageDataGenerator


# Load the CIFAR-10 dataset

(x_train, y_train), (x_test, y_test) = cifar10.load_data()


# Normalize the pixel values

x_train = x_train.astype('float32') / 255.0

x_test = x_test.astype('float32') / 255.0


# One-hot encode the labels

y_train = to_categorical(y_train, 10)

y_test = to_categorical(y_test, 10)


# Load the MobileNetV2 model pre-trained on ImageNet, excluding the top fully connected layers

base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(32, 32, 3))


# Freeze the layers of the base model

for layer in base_model.layers:

    layer.trainable = False


# Define the model

model = Sequential([

    base_model,

    GlobalAveragePooling2D(),
```



```

Dense(256, activation='relu'),
Dropout(0.5),
Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Data augmentation
datagen = ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True
)
datagen.fit(x_train)

# Train the model
model.fit(datagen.flow(x_train, y_train, batch_size=32), epochs=20, validation_data=(x_test, y_test))

# Evaluate the model
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test accuracy: {test_acc:.4f}")

```

## 9. Perform Sentiment Analysis using RNN

```

import tensorflow as tf

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences

# Load the IMDB dataset

```

```
max_features = 10000 # Only consider the top 10,000 words in the dataset
max_len = 500 # Only consider the first 500 words of each review
```

```
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)
```

```
# Pad sequences to ensure they are all the same length
```

```
x_train = pad_sequences(x_train, maxlen=max_len)
```

```
x_test = pad_sequences(x_test, maxlen=max_len)
```

```
# Define the model
```

```
model = Sequential([
    Embedding(max_features, 128, input_length=max_len),
    SimpleRNN(128),
    Dense(1, activation='sigmoid')
])
```

```
# Compile the model
```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
# Train the model
```

```
model.fit(x_train, y_train, epochs=10, batch_size=32, validation_split=0.2)
```

```
# Evaluate the model
```

```
test_loss, test_acc = model.evaluate(x_test, y_test)
```

```
print(f"Test accuracy: {test_acc:.4f}")
```

## **10. . Implement an LSTM based Autoencoder in TensorFlow/Keras.**

```
import tensorflow as tf
```

```
from tensorflow.keras.models import Model
```

```
from tensorflow.keras.layers import Input, LSTM, RepeatVector, Dense
```

```
from tensorflow.keras.datasets import imdb
```

```
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```

# Load the IMDB dataset

max_features = 10000 # Limit the number of words in the vocabulary

max_len = 100 # Maximum length of the input sequences

(x_train, _), (x_test, _) = imdb.load_data(num_words=max_features)

# Pad sequences to ensure uniform input size

x_train = pad_sequences(x_train, maxlen=max_len)

x_test = pad_sequences(x_test, maxlen=max_len)

# Define the LSTM Autoencoder model

timesteps = max_len

input_dim = max_features

# Encoder

encoder_inputs = Input(shape=(timesteps, input_dim))

encoded = LSTM(64, activation='relu', return_sequences=False)(encoder_inputs)

latent_dim = 32

encoded_latent = Dense(latent_dim, activation='relu')(encoded)

# Latent Space Representation

latent_inputs = Input(shape=(latent_dim,))

decoded_latent = Dense(64, activation='relu')(latent_inputs)

decoded = RepeatVector(timesteps)(decoded_latent)

decoded = LSTM(input_dim, activation='sigmoid', return_sequences=True)(decoded)

# Define the full Autoencoder model

autoencoder = Model(encoder_inputs, decoded)

# Define the Encoder model

encoder_model = Model(encoder_inputs, encoded_latent)

# Define the Decoder model

decoder_inputs = Input(shape=(latent_dim,))

decoder_outputs = autoencoder.layers[-2](decoder_inputs)

decoder_outputs = autoencoder.layers[-1](decoder_outputs)

decoder_model = Model(decoder_inputs, decoder_outputs)

```

```

# Compile the Autoencoder model

autoencoder.compile(optimizer='adam', loss='binary_crossentropy')


# Train the model

autoencoder.fit(x_train, x_train, epochs=10, batch_size=64, validation_split=0.2)


# Evaluate the model

loss = autoencoder.evaluate(x_test, x_test)

print(f"Test loss: {loss:.4f}")

```

## 11. . Image generation using GAN

```

import tensorflow as tf

from tensorflow.keras.datasets import mnist

from tensorflow.keras.layers import Dense, Reshape, Flatten, LeakyReLU, BatchNormalization,
Conv2DTranspose, Conv2D

from tensorflow.keras.models import Sequential

from tensorflow.keras.optimizers import Adam

import numpy as np

import matplotlib.pyplot as plt


# Load and preprocess the MNIST dataset

(x_train, _), (_, _) = mnist.load_data()

x_train = x_train / 255.0

x_train = np.expand_dims(x_train, axis=-1)


# Define the Generator model

def build_generator():

    model = Sequential()

    model.add(Dense(128 * 7 * 7, input_dim=100))

    model.add(LeakyReLU(alpha=0.2))

    model.add(Reshape((7, 7, 128)))

    model.add(BatchNormalization())

```

```
model.add(Conv2DTranspose(128, kernel_size=3, strides=2, padding='same'))
model.add(LeakyReLU(alpha=0.2))

model.add(Conv2DTranspose(64, kernel_size=3, strides=2, padding='same'))
model.add(LeakyReLU(alpha=0.2))

model.add(Conv2DTranspose(1, kernel_size=3, activation='sigmoid', padding='same'))

model.compile(loss='binary_crossentropy', optimizer=Adam())
return model
```

# Define the Discriminator model

```
def build_discriminator():
    model = Sequential()
    model.add(Conv2D(64, kernel_size=3, strides=2, padding='same', input_shape=(28, 28, 1)))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Conv2D(128, kernel_size=3, strides=2, padding='same'))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Flatten())
    model.add(Dense(1, activation='sigmoid'))

    model.compile(loss='binary_crossentropy', optimizer=Adam())
    return model
```

# Define the GAN model

```
def build_gan(generator, discriminator):
    model = Sequential()
    model.add(generator)
```

```
model.add(discriminator)

return model
```

```
# Create and compile models
```

```
generator = build_generator()
discriminator = build_discriminator()
discriminator.trainable = False
```

```
gan = build_gan(generator, discriminator)
gan.compile(loss='binary_crossentropy', optimizer=Adam())
```

```
# Training parameters
```

```
epochs = 10000
batch_size = 64
half_batch = batch_size // 2
```

```
# Training loop
```

```
for epoch in range(epochs):
```

```
    # Train Discriminator
```

```
    idx = np.random.randint(0, x_train.shape[0], half_batch)
    real_imgs = x_train[idx]
    real_labels = np.ones((half_batch, 1))
```

```
    noise = np.random.randn(half_batch, 100)
    fake_imgs = generator.predict(noise)
    fake_labels = np.zeros((half_batch, 1))
```

```
    d_loss_real = discriminator.train_on_batch(real_imgs, real_labels)
    d_loss_fake = discriminator.train_on_batch(fake_imgs, fake_labels)
    d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
```

```

# Train Generator

noise = np.random.randn(batch_size, 100)
valid_labels = np.ones((batch_size, 1))
g_loss = gan.train_on_batch(noise, valid_labels)


# Print progress

if epoch % 1000 == 0:

    print(f"{epoch}/{epochs} [D loss: {d_loss[0]} | D accuracy: {100 * d_loss[1]}] [G loss: {g_loss}]")


# Save generated images

if epoch % 1000 == 0:

    generated_images = generator.predict(np.random.randn(25, 100))

    generated_images = 0.5 * generated_images + 0.5 # Rescale images to [0, 1]

    fig, axs = plt.subplots(5, 5)

    cnt = 0

    for i in range(5):

        for j in range(5):

            axs[i,j].imshow(generated_images[cnt, :, :, 0], cmap='gray')

            axs[i,j].axis('off')

            cnt += 1

    plt.show()

```

## 12. Train a Deep learning model to classify a given image using pre trained model

```

import tensorflow as tf

from tensorflow.keras.applications import VGG16

from tensorflow.keras.preprocessing import image

from tensorflow.keras.applications.vgg16 import preprocess_input, decode_predictions

import numpy as np


# Define the path to your image

img_path = 'path_to_your_image.jpg' # Replace with the path to your image

```

```

# Load and preprocess the image

img = image.load_img(img_path, target_size=(224, 224)) # VGG16 expects 224x224 images
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Load the VGG16 model pre-trained on ImageNet
model = VGG16(weights='imagenet')

# Predict the class of the image
predictions = model.predict(img_array)

# Decode the predictions
decoded_predictions = decode_predictions(predictions, top=3)[0] # Get top 3 predictions

# Print the results
for i, (class_id, class_name, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {class_name} ({score:.2f})")

```

### **13. Recommendation system from sales data using Deep Learning**

```

import numpy as np
import pandas as pd

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Embedding, Flatten
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

# Example dataset
data = {
    'user_id': [1, 1, 2, 2, 3, 3, 4, 4],

```



```
'item_id': [1, 2, 1, 3, 2, 3, 1, 2],  
'rating': [5, 3, 4, 2, 5, 1, 2, 3]  
}
```

```
df = pd.DataFrame(data)
```

```
# Encode user and item IDs
```

```
user_encoder = LabelEncoder()
```

```
item_encoder = LabelEncoder()
```

```
df['user_id_encoded'] = user_encoder.fit_transform(df['user_id'])
```

```
df['item_id_encoded'] = item_encoder.fit_transform(df['item_id'])
```

```
# Define model parameters
```

```
n_users = len(user_encoder.classes_)
```

```
n_items = len(item_encoder.classes_)
```

```
embedding_dim = 10
```

```
# Build the recommendation model
```

```
model = Sequential([
```

```
    Embedding(input_dim=n_users, output_dim=embedding_dim, input_length=1,  
    name='user_embedding'),
```

```
    Flatten(),
```

```
    Dense(embedding_dim, activation='relu'),
```

```
    Dense(n_items, activation='sigmoid') # Use sigmoid for rating prediction
```

```
])
```

```
model.compile(optimizer=Adam(), loss='binary_crossentropy') # Use 'binary_crossentropy' for  
simplicity
```

```
# Prepare input features and target
```

```
X = df[['user_id_encoded', 'item_id_encoded']].values
```

```
y = df['rating'].values
```

```
# Train/test split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Train the model
```

```
model.fit([X_train[:, 0], X_train[:, 1]], y_train, epochs=10, batch_size=2, validation_split=0.2)
```

```
# Evaluate the model
```

```
loss = model.evaluate([X_test[:, 0], X_test[:, 1]], y_test)
```

```
print(f"Test loss: {loss:.4f}")
```

#### **14. Implement Object Detection using CNN**

```
import tensorflow as tf
```

```
import tensorflow_hub as hub
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import matplotlib.patches as patches
```

```
from PIL import Image
```

```
# Load the pre-trained object detection model from TensorFlow Hub
```

```
model_url = "https://tfhub.dev/google/faster_rcnn/openimages_v4/inception_resnet_v2/1"
```

```
detector = hub.load(model_url).signatures['default']
```

```
# Load and preprocess the image
```

```
def load_image(image_path):
```

```
    img = Image.open(image_path)
```

```
    img = img.convert('RGB') # Ensure image is in RGB format
```

```
    img_np = np.array(img)
```

```
    return img_np
```

```
def preprocess_image(image_np):
```

```
image_tensor = tf.convert_to_tensor(image_np, dtype=tf.float32)
image_tensor = tf.image.convert_image_dtype(image_tensor, dtype=tf.uint8)
image_tensor = tf.expand_dims(image_tensor, 0) # Add batch dimension
return image_tensor
```

*# Perform object detection on the input image*

```
def detect_objects(image_tensor):
    result = detector(image_tensor)
    return result
```

```
def draw_boxes(image_np, boxes, class_ids, scores):
    fig, ax = plt.subplots(1, figsize=(12, 9))
    ax.imshow(image_np)
```

```
    for i in range(len(boxes)):
        box = boxes[i]
        class_id = class_ids[i]
        score = scores[i]
```

*# Draw bounding box*

```
    ymin, xmin, ymax, xmax = box
    rect = patches.Rectangle((xmin * image_np.shape[1], ymin * image_np.shape[0]),
                             (xmax - xmin) * image_np.shape[1],
                             (ymax - ymin) * image_np.shape[0],
                             linewidth=2, edgecolor='r', facecolor='none')
    ax.add_patch(rect)
    plt.text(xmin * image_np.shape[1], ymin * image_np.shape[0],
             f'{class_id} {{score:.2f}}', color='red', fontsize=12, bbox=dict(facecolor='yellow', alpha=0.5))

plt.show()
```

*# Example usage*

image\_path = 'path\_to\_your\_image.jpg' *# Replace with the path to your image*

image\_np = load\_image(image\_path)

image\_tensor = preprocess\_image(image\_np)

result = detect\_objects(image\_tensor)

*# Extract and draw results*

boxes = result['detection\_boxes'].numpy()[0]

class\_ids = result['detection\_classes'].numpy()[0]

scores = result['detection\_scores'].numpy()[0]

draw\_boxes(image\_np, boxes, class\_ids, scores)

## **15. Implement any simple Reinforcement Algorithm for an NLP problem**

import numpy as np

import tensorflow as tf

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dense, Embedding, LSTM

from tensorflow.keras.optimizers import Adam

import nltk

from nltk.tokenize import word\_tokenize

nltk.download('punkt')

class TextGenerationEnv:

def \_\_init\_\_(self, vocab\_size, max\_sequence\_length):

self.vocab\_size = vocab\_size

self.max\_sequence\_length = max\_sequence\_length

self.current\_step = 0

self.text\_sequence = np.zeros(max\_sequence\_length, dtype=int)

def reset(self):

self.current\_step = 0

```
self.text_sequence = np.zeros(self.max_sequence_length, dtype=int)

return self.text_sequence
```

```
def step(self, action):

    self.text_sequence[self.current_step] = action

    reward = self.compute_reward(self.text_sequence)

    self.current_step += 1

    done = self.current_step >= self.max_sequence_length

    return self.text_sequence, reward, done
```

```
def compute_reward(self, text_sequence):

    # For simplicity, we use the length of the sequence as a reward

    return np.sum(text_sequence)
```

```
def build_policy_network(vocab_size, embedding_dim, hidden_units):

    model = Sequential([

        Embedding(input_dim=vocab_size, output_dim=embedding_dim, input_length=1),

        LSTM(hidden_units, return_sequences=False),

        Dense(vocab_size, activation='softmax')

    ])

    return model
```

```
def train_policy_network(env, policy_network, optimizer, episodes=1000):

    for episode in range(episodes):

        state = env.reset()

        done = False

        total_reward = 0

        while not done:

            state = np.expand_dims(state, axis=0) # Add batch dimension

            action_probs = policy_network.predict(state)

            action = np.argmax(action_probs)
```

```
next_state, reward, done = env.step(action)
```

```
total_reward += reward
```

```
# Calculate the loss and update the policy network
```

```
with tf.GradientTape() as tape:
```

```
    action_probs = policy_network(state, training=True)
```

```
    loss = -tf.math.log(action_probs[0, action]) * reward
```

```
    grads = tape.gradient(loss, policy_network.trainable_variables)
```

```
    optimizer.apply_gradients(zip(grads, policy_network.trainable_variables))
```

```
state = next_state
```

```
print(f'Episode {episode+1}, Total Reward: {total_reward}')
```

```
# Parameters
```

```
vocab_size = 10 # Example vocabulary size
```

```
embedding_dim = 8
```

```
hidden_units = 16
```

```
max_sequence_length = 10
```

```
learning_rate = 0.001
```

```
# Initialize environment and policy network
```

```
env = TextGenerationEnv(vocab_size, max_sequence_length)
```

```
policy_network = build_policy_network(vocab_size, embedding_dim, hidden_units)
```

```
optimizer = Adam(learning_rate)
```

```
# Train policy network
```

```
train_policy_network(env, policy_network, optimizer)
```

